

8/30/2016 CS535 Big Data - Fall 2016 W2.A.1

CS535 BIG DATA

PART 0. INTRODUCTION
3. DATA MODEL FOR BIG DATA
: APACHE THRIFT

Sangmi Lee Pallickara
Computer Science, Colorado State University
<http://www.cs.colostate.edu/~cs535>

8/30/2016 CS535 Big Data - Fall 2016 W2.A.1

FAQs

- Programming Assignment 1
 - Due Sept. 28
 - Submission via Canvas
 - Please check the course Web Page at least twice a week

8/30/2016 CS535 Big Data - Fall 2016 W2.A.2

Today's topics

- Apache Thrift continued
- Introduction to MapReduce

8/30/2016 CS535 Big Data - Fall 2016 W2.A.3

Thrift Architecture

8/30/2016 CS535 Big Data - Fall 2016 W2.A.4

Protocol

8/30/2016 CS535 Big Data - Fall 2016 W2.A.5

Protocols (1/2)

- Describes "what" is transmitted
- Thrift supports both text and binary protocols
 - The binary protocol outperforms text protocol
 - The text protocol may be useful (such as in debugging)
- TBinaryProtocol
 - A straight-forward binary format encoding numeric values as binary, rather than converting to text.
- TCompactProtocol
 - Very efficient, dense encoding of data

8/30/2016 CS535 Big Data - Fall 2016 W2.A.6


Protocols (1/2)

- **TDenseProtocol**
 - Similar to **TCompactProtocol** but strips off the meta information from what is transmitted, and adds it back in at the receiver.
- **TJSONProtocol**
 - Uses JSON for encoding of data.
- **TSimpleJSONProtocol**
 - A write-only protocol using JSON. It cannot be parsed by Thrift. Suitable for parsing by scripting languages
- **TDebugProtocol**
 - Uses a human-readable text format to aid in debugging.

8/30/2016 CS535 Big Data - Fall 2016 W2.A.7

TCompactProtocol

- The **TCompactProtocol** is the most-efficient method in the Java implementation of Thrift
 - Writes numeric tags for each piece of data
 - The recipient is expected to properly match these tags with the data
 - If the data is not present, there simply is no tag/data pair



8/30/2016 CS535 Big Data - Fall 2016 W2.A.8

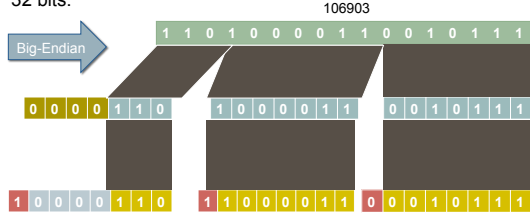
TCompactProtocol (1/2)

- For integers, the **TCompactProtocol** performs compression Variable-Length Quantity (VLQ) encoding
- VLQ uses 7 of 8 bits out of each byte for information
 - the 8th bit used as a continuation bit

8/30/2016 CS535 Big Data - Fall 2016 W2.A.9

TCompactProtocol (2/2)

- Decimal value 106903 (0x1A197) saving 1 byte if it was stored in 32 bits:



VLQ's worst-case encoding
For a 32-bit int, the worst case is 5 bytes
For 64-bit ints, the worst case is 10 bytes

8/30/2016 CS535 Big Data - Fall 2016 W2.A.10

Versioning

8/30/2016 CS535 Big Data - Fall 2016 W2.A.11

Versioning

- **Applications** evolve over time
 - You added an extra field to your message format
- The system **must be able to support reading of old data** from log files
 - Process requests from out-of-date clients to new servers and vice versa

8/30/2016 CS535 Big Data - Fall 2016 W2.A.12

Versioning Field Identifiers

- Any new fields should be optional
 - Any messages serialized by code using the "old" message format **can be parsed** by the new code
 - Without any missing required fields (Parsing error)
 - Messages created by new code can be parsed by old code
 - Old binary will ignore the new field when parsing
 - Unknown fields are not discarded: it will be serialized along with other fields
- Non-required fields can be removed
- Changing a default value is allowed

Mark Slee, Aditya Agarwal and Marc Kwiatkowski, Thrift: Scalable Cross-Language Services Implementation, <https://thrift.apache.org/static/files/thrift-20070401.pdf>

CS535 BIG DATA

PART 1. BATCH COMPUTING MODELS FOR BIG DATA ANALYTICS

1. DISTRIBUTED MODEL FOR SCALABLE BATCH COMPUTING - MapReduce

Sangmi Lee Pallickara
Computer Science, Colorado State University
<http://www.cs.colostate.edu/~cs535>

8/30/2016 CS535 Big Data - Fall 2016 W2.A.14

This material is built based on

- Jeffrey Dean and Sanjay Ghemawat, "MapReduce: Simplified Data Processing on Large Clusters" In Proceeding OSDI'04 Proceedings of the 6th conference on Symposium on Operating Systems Design & Implementation - Vol. 6
- Hadoop: The definitive Guide, Tom White, O'Reilly, 3rd Edition, 2014
- MapReduce Design Patterns, Donald Miner and Adam Shook, O'Reilly, 2013
- Anand Rajaraman, Jure Leskovec, and Jeffrey Ullman, "Mining of Massive Datasets", Cambridge University Press, 2012 -- Chapter 2

8/30/2016 CS535 Big Data - Fall 2016 W2.A.15

Why MapReduce?

8/30/2016 CS535 Big Data - Fall 2016 W2.A.16

What is MapReduce?

- MapReduce is inspired by the concepts of *map* and *reduce* in Lisp.

8/30/2016 CS535 Big Data - Fall 2016 W2.A.17

Python implementations

Python implementations

- Cython
- Brython
- CLPython
- HotPy
- IronPython
- Jython
- Pyjs
- PyPy
- SNAPPY ...

```

# Grammar for Python
# Note: Changing the grammar specified in this file will most likely
# require corresponding changes in the parser module
# (.Modules parameter module c). If you can't make the changes to
# that module yourself, please co-ordinate the required changes
# with someone who can; ask around on python-dev for help. Fred
# Drake <drake@bcom.org> will probably be listening there.

# NOTE WELL: You should also follow all the steps listed at
# https://docs.python.org/devguide/grammar.html

# Start symbols for the grammar:
# single_input is a single interactive statement;
# file_input is a module or sequence of commands read from an input file;
# eval_input is the input for the eval() function.
# NB: compound_stmt in single_input is followed by extra NEWLINE
single_input: NEWLINE | single_stmt | compound_stmt NEWLINE
file_input: (NEWLINE | stmt)* ENDMARKER
eval_input: (NEWLINE | stmt)* ENDMARKER

# For normal assignments, additional restrictions enforced by the
# interpreter
del_stmt: 'del' exprlist
pass_stmt: 'pass'
flow_stmt: break_stmt | continue_stmt | return_stmt | raise_stmt |
yield_stmt
break_stmt: 'break'
continue_stmt: 'continue'
return_stmt: 'return' (testlist)
yield_stmt: 'yield' expr
raise_stmt: 'raise' (test ['from' test])
import_stmt: import_name | import_from
import_name: 'import' dotted_as_names
# note below: the ('.' '.') is necessary because '.' is tokenized as
ELLIPSIS
import_from: ('from' (('.' '.') dotted_name | ('.' '.')*)
import_as_name: NAME ['as' NAME]
dotted_as_name: dotted_name ['as' NAME]
import_as_names: import_as_name (',' import_as_name)*
dotted_as_names: dotted_as_name (',' dotted_as_name)*
global_stmt: 'global' NAME (',' NAME)*
nonlocal_stmt: 'nonlocal' NAME (',' NAME)*
assert_stmt: 'assert' test [',' test]
  
```

Programming language specification: Python 3.5.2

8/30/2016 CS535 Big Data - Fall 2016 W2.A.18

What is MapReduce?

- Developed within Google as a mechanism for processing large amounts of raw data.
 - Crawled documents or web request logs
 - Distributes these data across thousands of machines
 - Same computations are performed on each CPU with a different portion of the dataset

8/30/2016 CS535 Big Data - Fall 2016 W2.A.19

Why MapReduce?

- MapReduce provides an abstraction that allows engineers to perform simple computations while hiding the details of:
 - Parallelization
 - Data distribution
 - Load balancing
 - Fault tolerance

8/30/2016 CS535 Big Data - Fall 2016 W2.A.20

Example

- Climate data from National Climatic Data Center (NCDC)

```
0057
332130 # USAF weather station identifier
99999 # WBAN weather station identifier
19500101 # observation date
0300 # observation time 4
+ 51317 # latitude (degrees x 1000)
+ 028783 # longitude (degrees x 1000)
FM-12
+ 0171 # elevation (meters)
99999
V020
320 # wind direction (degrees)
1 # quality code
```

8/30/2016 CS535 Big Data - Fall 2016 W2.A.21

The first entries for 1990

```
% ls raw/ 1990 | head
010010-99999-1990.gz
010014-99999-1990.gz
010015-99999-1990.gz
010016-99999-1990.gz
010017-99999-1990.gz
010030-99999-1990.gz
010040-99999-1990.gz
010080-99999-1990.gz
010100-99999-1990.gz
010150-99999-1990.gz
```

8/30/2016 CS535 Big Data - Fall 2016 W2.A.22

Analyzing the data with Unix Tools (1/2)

- A program for finding the **maximum recorded temperature by year** from NCDC weather records
 - e.g. Weather change for a century

```
#!/usr/bin/env bash
for year in $(ls */* | sed 's/^[^/]*/' | sort | uniq); do
    echo -ne "basename $year .gz " `t`
    gunzip -c $year | \
    awk '{
        temp = substr($0, 88, 5) + 0;
        q = substr($0, 93, 1);
        if (temp != 9999 && q ~ /[01459]/ && temp > max)
            max = temp;
    }
    END { print max }'
done
```

8/30/2016 CS535 Big Data - Fall 2016 W2.A.23

Analyzing the data with Unix Tools (2/2)

- The script loops through the compressed year files
 - Printing the year
 - Processing each file using awk
 - Extracts two fields
 - Air temperature and the quality code
 - Check if it is greater than the maximum value seen so far

```
% ./max_temperature.sh
1901 317
1902 244
1903 289
1904 256
1905 283
...
```

8/30/2016 CS535 Big Data - Fall 2016 W2.A.24

Results?

- The complete run for the century took 42 minutes
- To speed up the processing
 - We need to run parts of the program in parallel
- Process different years in different processes
- What will be the problems?

8/30/2016 CS535 Big Data - Fall 2016 W2.A.25

Challenges

- Dividing the work** into equal-size pieces
 - Data size per year?
- Combining the results** from independent processes
 - Combining results and sorting by year?

8/30/2016 CS535 Big Data - Fall 2016 W2.A.26

Map and Reduce

- MapReduce works by breaking the processing into two phases
 - The map phase
 - The reduce phase
- Each phase has key-value pairs as input and output
- Programmers should specify
 - Types of input/output key-values
 - The map function
 - The reduce function

8/30/2016 CS535 Big Data - Fall 2016 W2.A.27

Visualizing the way MapReduce works (1/4)

Sample lines of input data

```
0067011990999991950051507004... 9999999N9 + 00001 +9999999999...
0043011990999991950051512004... 9999999N9 + 00221 +9999999999...
0043011990999991950051518004... 9999999N9-00111 +9999999999...
0043012650999991949032412004... 0500001N9 + 01111 +9999999999...
0043012650999991949032418004... 0500001N9 + 00781 +9999999999...
```

These lines are presented to the map function as key-value pairs

```
(0, 0067011990999991950051507004...9999999N9 + 00001+
9999999999...)
(106, 0043011990999991950051512004...9999999N9 + 00221+
9999999999...)
(212, 0043011990999991950051518004...9999999N9-0011 1 +
9999999999...)
```

The keys are the line offsets within the file

8/30/2016 CS535 Big Data - Fall 2016 W2.A.28

Visualizing the way MapReduce works (2/4)

The **map function** extracts **the year and the air temperature** and emits them as its output

```
(1950, 0)
(1950, 22)
(1950, -11)
(1949, 111)
(1949, 78)
```

This output key-value pairs will be **sorted and grouped by key**. Our reduce function will see the following input:

```
(1949, [111, 78])
(1950, [0, 22, -11])
```

8/30/2016 CS535 Big Data - Fall 2016 W2.A.29

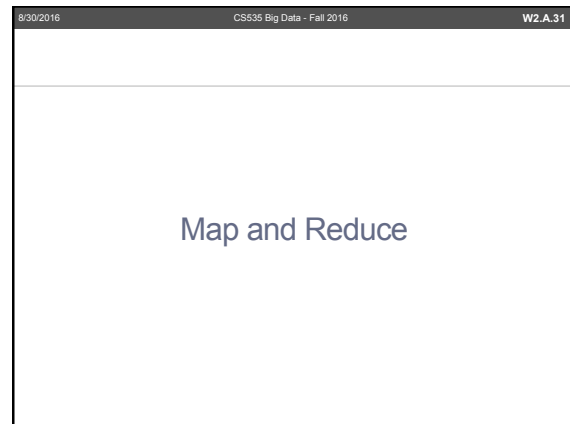
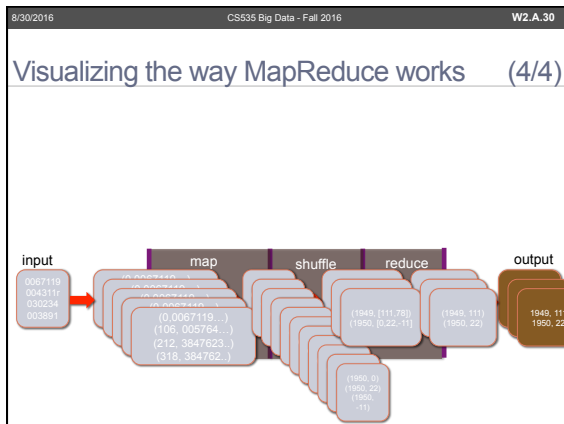
Visualizing the way MapReduce works (3/4)

Reduce function iterates through the list and picks the maximum reading

```
(1949, 111)
(1950, 22)
```

This is the final output

```
graph LR
    Input["input  
(0067119...)  
(0043119...)  
(000224...)  
(003891...)"] --> Map["map  
(0, 0067119...)  
(106, 005764...)  
(212, 3847623...)  
(318, 384762...)"]
    Map --> Shuffle["shuffle  
(1950, 0)  
(1950, 22)  
(1950, -11)"]
    Shuffle --> Reduce["reduce  
(1949, [111, 78])  
(1950, [0, 22, -11])"]
    Reduce --> Output["output  
1949, 111  
1950, 22"]
```



8/30/2016 CS535 Big Data - Fall 2016 W2.A.32

Map function

- Takes an input pair
- **Produces a set of intermediate key/value pairs**
- The MapReduce library groups together all intermediate values associated with the same intermediate key i and passes them to the reduce function

8/30/2016 CS535 Big Data - Fall 2016 W2.A.33

Reduce function

- **Accepts an intermediate key i** and a set of values for that key
- **Merges** together these values
 - Forms a possibly smaller **set** of values

8/30/2016 CS535 Big Data - Fall 2016 W2.A.34

Example: counting words

- Count the number of occurrences of each word in a large collection of documents:

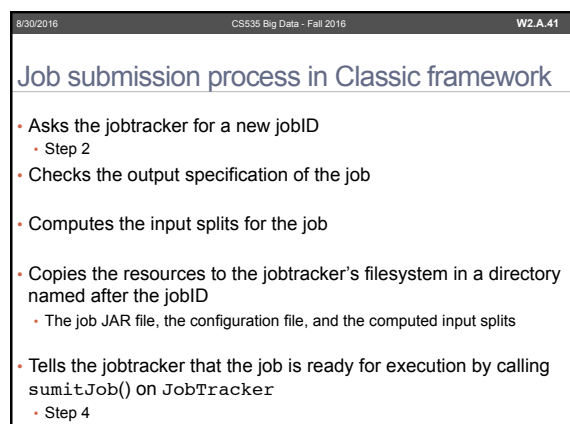
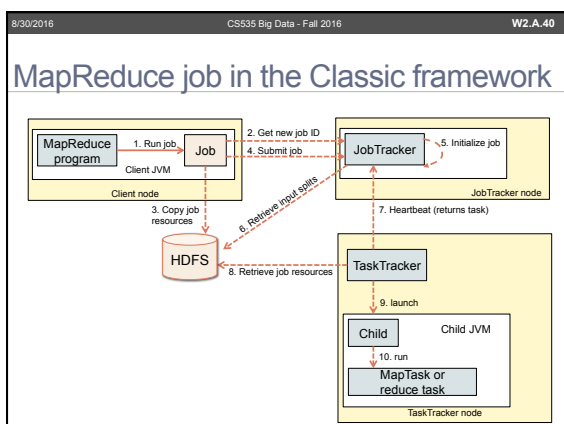
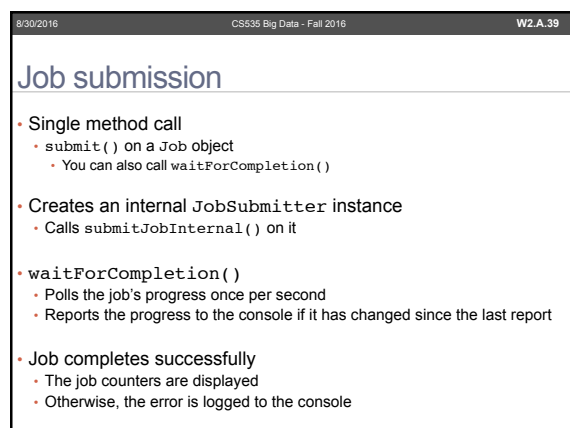
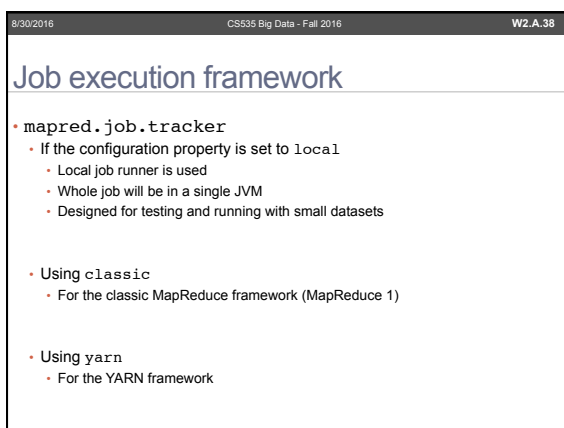
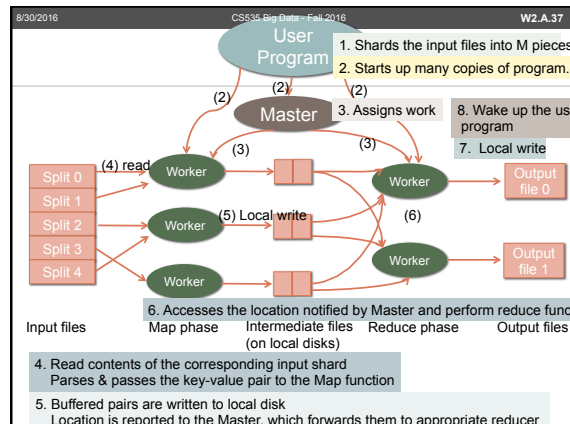
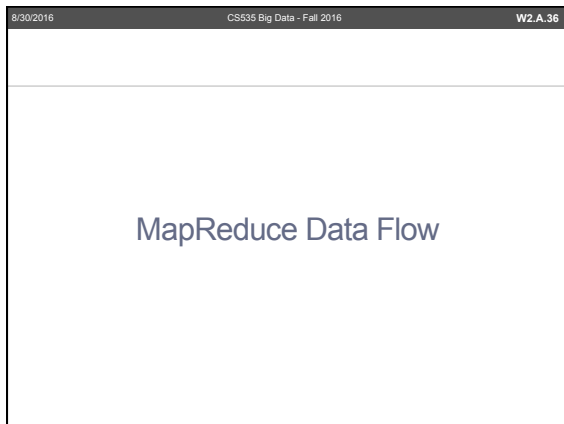
```
map(String key, String value):
// key: document name
// value: document contents
for each word w in value:
    EmitIntermediate(w, "1");

reduce(String key, Iterator values):
//key: a word
//values: a list of counts
int result = 0;
for each v in values:
    result += ParseInt(v)
Emit(AsString(result));
```

8/30/2016 CS535 Big Data - Fall 2016 W2.A.35

Comparison with other systems

- **MPI vs. MapReduce**
 - MapReduce tries to collocate the data with the compute node
 - Data access is fast
 - Data is local!
- **Volunteer computing vs. MapReduce**
 - SETI@home
 - Uses donated (volunteered) CPU time



8/30/2016 CS535 Big Data - Fall 2016 W2.A.42

Job Initialization in the Classic mode

- JobTracker puts the job into an internal queue
 - Creating an object to represent the job being run
 - Bookkeeping information to keep track of the status and progress of its tasks
- Creating the list of tasks
 - The job scheduler retrieves the input splits
 - The job scheduler creates a map task for each split
 - The number of reduce tasks to create
 - `mapred.reduce.tasks` property in the Job
 - `setNumReduceTasks()`
 - A job setup task and a job cleanup task are created

8/30/2016 CS535 Big Data - Fall 2016 W2.A.43

Task assignment

- Jobtracker select a job to run
 - Default algorithm
 - Based on the priority list of jobs
- Tasktrackers have a fixed number of slots for map tasks and reduce tasks
 - These are set independently
 - These are selected based on the number of cores and the memory
- The default scheduler fills empty map task slots first
 - Before it fills the reduce task slots
 - Scheduling the reduce task does not need to consider data locality
- Tasktrackers run a simple loop that periodically sends heartbeat method calls to the jobtracker

8/30/2016 CS535 Big Data - Fall 2016 W2.A.44

Task execution

- Tasktracker
 - Copies the job Jar to the tasktracker's file system
 - Any files needed from the distributed cache to the local file system
 - Creates a local working directory
 - Creates instance of TaskRunner
- TaskRunner
 - Launches a new Java Virtual Machine to run each task
 - Any failed map or reduce does not affect the tasktracker

8/30/2016 CS535 Big Data - Fall 2016 W2.A.45

Streaming and pipes

- Runs special map and reduce tasks
 - To launch the user supplied executable
 - To communicate with it
- Streaming task communicates with the process using standard input and output streams
- Pipes task listens on a socket and passes the C++ process a port number in its environment
 - On startup, the C++ process establish a persistent socket connection back to the parent Java Pipes task

8/30/2016 CS535 Big Data - Fall 2016 W2.A.46

Progress and status updates [1/3]

- A Job and each of its tasks have a status
 - State of the job or task
 - E.g. running, successfully complete, failed
 - The progress of maps and reduces
 - The values of the job's counters
 - A status message or description
- Progress of a task
 - The proportion of the task completed
 - Map task
 - The proportion of the input that has been processed
 - Reduce task
 - Divides the total progress into 3 parts (copy/sort/reduce)
 - If the task has run the reducer on half its input
 - $1/3 \text{ (copy)} + 1/3 \text{ (sort)} + \text{a half of } 1/3 \text{ (reduce phase)} = 5/6$

8/30/2016 CS535 Big Data - Fall 2016 W2.A.47

Progress and status updates [2/3]

- Tasks
 - have a set of counters
 - Count various events at the task run
 - E.g. the number of map output records written
 - If a task reports progress
 - it sets a flag to indicate that the status change should be sent to the tasktracker
 - Checked every 3 seconds
- Tasktracker
 - Tasks notify the current task status to the tasktracker
 - if the flag is set
 - Tasktracker sends heartbeats to the jobtracker every 5 seconds (minimum)

8/30/2016 CS535 Big Data - Fall 2016 W2.A.48

Progress and status updates [2/3]

- Jobtracker
 - Combines updates to produce a global view
- Job
 - Receives the latest status by polling the jobtracker every second
 - Prints job statistics and counters to the console

8/30/2016 CS535 Big Data - Fall 2016 W2.A.49

YARN (MapReduce 2)

- To provide the scalability to MapReduce
 - Splitting responsibility of the jobtracker
 - Scheduling
 - Task progress monitoring
- MapReduce is one type of YARN application