

9/8/2016 CS535 Big Data - Fall 2016 W3.B.0

## Computer Science Department Picnic

**Welcome to the 2016-2017 Academic year !**

Meet your faculty, department staff, and fellow students in a social setting. Food and drink will be provided.



**When: Saturday, September 10<sup>th</sup>**  
**Time: 11am – 2pm**  
**Where: City Park Shelter #7**

CS535 BIG DATA

---

**PART 1. BATCH COMPUTING MODELS FOR BIG DATA ANALYTICS**  
**1. DISTRIBUTED MODEL FOR SCALABLE BATCH COMPUTING - MapReduce**

---

Sangmi Lee Pallickara  
Computer Science, Colorado State University  
<http://www.cs.colostate.edu/~cs535>

9/8/2016 CS535 Big Data - Fall 2016 W3.B.2

## FAQs

- Questions about PA1
  - Send an email to [cs535@cs.colostate.edu](mailto:cs535@cs.colostate.edu)
- Use the posted configuration file with 2GB memory for the worker node

9/8/2016 CS535 Big Data - Fall 2016 W3.B.3

## Objectives

- In-Memory Cluster Computing
- Introduction to Apache Spark
- RDD
- Spark cluster
- Scheduling

9/8/2016 CS535 Big Data - Fall 2016 W3.B.4

## In-Memory Cluster Computing: Apache Spark

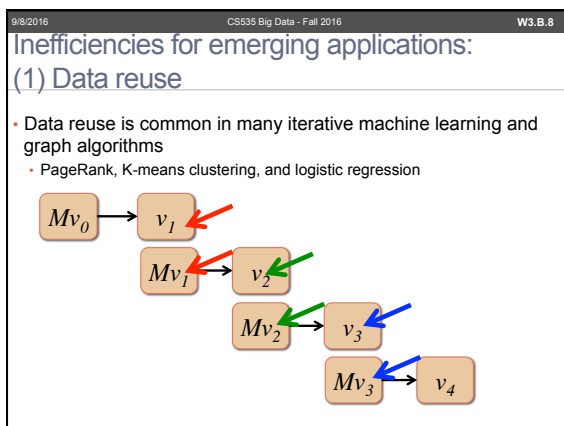
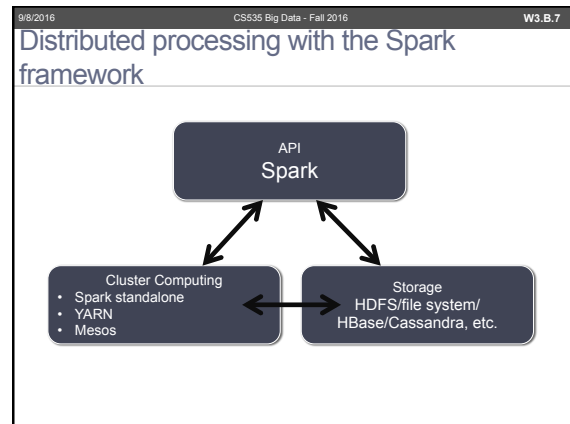
9/8/2016 CS535 Big Data - Fall 2016 W3.B.5

## In-Memory Cluster Computing: Apache Spark Introduction

9/8/2016 CS535 Big Data - Fall 2016 W3.B.6

### This material is built based on

- Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, and Ion Stoica, "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing," The 9th USENIX Symposium on Networked Systems Design and Implementation (NSDI 12)
- Holden Karau, Andy Komwinski, Patrick Wendell and Matei Zaharia, "Learning Spark", O'Reilly, 2015
- Spark Overview, <https://spark.apache.org/docs/2.0.0-preview/>
  - Spark programming guide  
<https://spark.apache.org/docs/2.0.0-preview/programming-guide.html>
  - Job Scheduling  
<https://spark.apache.org/docs/2.0.0-preview/job-scheduling.html>



9/8/2016 CS535 Big Data - Fall 2016 W3.B.9

### Inefficiencies for emerging applications: (2) Interactive data analytics

- User runs multiple ad-hoc queries on the **same subset** of the data

9/8/2016 CS535 Big Data - Fall 2016 W3.B.10

### Existing approaches

- Hadoop
  - Writing output to an external stable storage system
    - e.g. HDFS
    - Substantial overheads due to data replication, disk I/O, and serialization
- Pregel
  - Iterative graph computations
- HaLoop
  - Iterative MapReduce interface
- Pregel/HaLoop support specific computation patterns
  - e.g. looping a series of MapReduce steps

9/8/2016 CS535 Big Data - Fall 2016 W3.B.11

### In-Memory Cluster Computing: Apache Spark RDD (Resilient Distributed Dataset)

9/8/2016 CS535 Big Data - Fall 2016 W3.B.12

## RDD (Resilient Distributed Dataset)

- **Read-only, partitioned collection of records**
  - A fault-tolerant collection of elements that can be operated on in parallel
- RDDs are the core unit of data in Spark
  - Most Spark programming involves performing operations on RDDs

9/8/2016 CS535 Big Data - Fall 2016 W3.B.13

## Overview of RDD

- **Lineage**
  - How it was derived from other dataset to compute its partitions from data in stable storage?
    - RDDs do not need to be materialized at all times
  - Program **CANNOT** reference an RDD if it cannot reconstruct after a failure
- **Persistence**
  - Users can indicate which RDDs they will reuse and the storage strategy
- **Partitioning**
  - Users can specify the partitioning method across machines based on a key in each record

9/8/2016 CS535 Big Data - Fall 2016 W3.B.14

## Spark Programming Interface to RDD [1/3]

- **transformations**
  - Operations that create RDDs
    - Return pointers to new RDDs
    - e.g. map, filter, and join
  - RDDs can only be created through deterministic operations on either
    - Data in stable storage
    - Other RDDs

9/8/2016 CS535 Big Data - Fall 2016 W3.B.15

## Spark Programming Interface to RDD [2/3]

- **actions**
  - Operations that return a value to the application or export data to a storage system
    - e.g. count: returns the number of elements in the dataset
    - e.g. collect: returns the elements themselves
    - e.g. save: outputs the dataset to a storage system

9/8/2016 CS535 Big Data - Fall 2016 W3.B.16

## Spark Programming Interface to RDD [3/3]

- **persist**
  - Indicates which RDDs they want to **reuse in future operations**
  - Spark keeps persistent RDDs **in memory** by default
  - If there is not enough RAM
    - It can spill them to disk
  - Users are allowed to,
    - store the RDD only on disk
    - replicate the RDD across machines
    - specify a persistence priority on each RDD

9/8/2016 CS535 Big Data - Fall 2016 W3.B.17

## Example: Console Log Mining [1/3]

- Suppose that a web service is experiencing errors and an operator wants to search terabytes of logs in the Hadoop file system (HDFS) to find the cause
- The user can load just **the error messages** from the logs into the RAM across a set of nodes and query them interactively

```
lines = spark.textFile("hdfs://...")
errors=lines.filter(_.startsWith("ERROR"))
errors.persist()
```

No work has been performed  
User can use the RDD in actions

9/8/2016 CS535 Big Data - Fall 2016 W3.B.18

## Example: Console Log Mining [2/3]

- Users can perform further transformations and actions on the RDD

```
//To count number of error messages
errors.count()

//Count errors mentioning MySQL:
errors.filter(_.contains("MySQL")).count()

//Return the time fields of errors mentioning
//HDFS as an array (assuming time is field
//number 3 in a tab-separated format
errors.filter(_.contains("HDFS"))
.map(_.split('/t')(3))
.collect()
```

After the first action involving errors runs, Spark will store the partitions of errors in memory.

9/8/2016 CS535 Big Data - Fall 2016 W3.B.19

## Example: Console Log Mining [3/3]

Spark code

```
lines = spark.textFile("hdfs://...")
errors=lines.filter(_.startsWith("
ERROR"))
errors.persist()
errors.filter(_.contains("HDFS"))
.map(_.split('/t')(3))
.collect()
```

Lineage graph

```
graph TD
    lines -->|filter(_.startsWith("ERROR"))| errors
    errors -->|filter(_.contains("HDFS"))| HDFS_errors[HDFS errors]
    HDFS_errors -->|map(_.split('/t')(3))| time_fields[Time fields]
```

If a partition of errors is lost Spark rebuilds it by applying a filter on only the corresponding partition of lines

9/8/2016 CS535 Big Data - Fall 2016 W3.B.20

## Lazy Evaluation

- Transformations on RDDs are **lazily evaluated**
  - Spark will **NOT** begin to execute until it sees an action
  - Spark internally records metadata to indicate that this operation has been requested
- Loading data into an RDD is lazily evaluated
- Reduces the number of passes it has to take over our data by grouping operations together

9/8/2016 CS535 Big Data - Fall 2016 W3.B.21

## Benefits of RDDs as a distributed memory abstraction [1/3]

- RDDs can only be created ("written") through **coarse-grained transformations**
  - Distributed shared memory (DSM) allows reads and writes to each memory location
  - Reads on RDDs can still be fine-grained
    - A large read-only lookup table
  - Applications perform bulk writes
  - More efficient fault tolerance
    - Lineage based bulk recovery

9/8/2016 CS535 Big Data - Fall 2016 W3.B.22

## Benefits of RDDs as a distributed memory abstraction [2/3]

- RDDs' immutable data
  - System can mitigate slow nodes (Stragglers)
    - Creates backup copies of slow tasks
      - without accessing the same memory
  - Spark distributes the data over different working nodes that run computations in parallel
    - Orchestrates communicating between nodes to integrate intermediate results and combine them for the final result

9/8/2016 CS535 Big Data - Fall 2016 W3.B.23

## Benefits of RDDs as a distributed memory abstraction [3/3]

- Runtime can schedule tasks based on data locality
  - To improve performance
- RDDs degrade gracefully when there is insufficient memory
  - Partitions that do not fit in the RAM are stored on disk

9/8/2016 CS535 Big Data - Fall 2016 W3.B.24

## Applications not suitable for RDDs

- RDDs are best suited for batch applications that apply the same operations to all elements of a dataset
  - Steps are managed by lineage graph efficiently
  - Recovery is managed effectively
- RDDs would not be suitable for applications
  - Making asynchronous fine-grained updates to shared state
  - e.g. a storage system for a web application or an incremental web crawler

9/8/2016 CS535 Big Data - Fall 2016 W3.B.25

## In-Memory Cluster Computing: Apache Spark RDDs in Spark

9/8/2016 CS535 Big Data - Fall 2016 W3.B.26

## RDDs in Spark: The Runtime

User's driver program launches multiple workers, which read data blocks from a distributed file system and can persist computed RDD partitions in memory

9/8/2016 CS535 Big Data - Fall 2016 W3.B.27

## Representing RDDs

- A set of partitions
  - Atomic pieces of the dataset
- A set of dependencies on parent RDDs
- A function for computing the dataset based on its parents
- Metadata about its partitioning scheme
- Data placement

9/8/2016 CS535 Big Data - Fall 2016 W3.B.28

## In-Memory Cluster Computing: Apache Spark RDD Dependency in Spark

9/8/2016 CS535 Big Data - Fall 2016 W3.B.29

## Dependency between RDDs [1/2]

- *Narrow* dependency
- *Wide* dependency

9/8/2016 CS535 Big Data - Fall 2016 W3.B.30

## Dependency between RDDs [1/2]

- Narrow dependency**
  - Each partition of the parent RDD is used by **at most one partition** of the child RDD

The diagram shows four examples of narrow dependencies:
 

- map, filter:** A 2x2 grid of parent partitions connected to a 2x2 grid of child partitions, with one-to-one vertical connections.
- union:** A 2x2 grid of parent partitions connected to a 4x2 grid of child partitions, with each parent partition connected to exactly one child partition.
- Join with inputs co-partitioned:** A 2x2 grid of parent partitions connected to a 2x2 grid of child partitions, with each parent partition connected to exactly one child partition.

9/8/2016 CS535 Big Data - Fall 2016 W3.B.31

## Dependency between RDDs [1/2]

- Wide dependency**
  - Multiple child partitions** may depend on a single partition of parent RDD

The diagram shows two examples of wide dependencies:
 

- groupByKey:** A 2x2 grid of parent partitions connected to a 2x2 grid of child partitions, with each parent partition connected to all child partitions in its column.
- Join with inputs not co-partitioned:** A 2x2 grid of parent partitions connected to a 2x2 grid of child partitions, with each parent partition connected to all child partitions.

9/8/2016 CS535 Big Data - Fall 2016 W3.B.32

## Dependency between RDDs [2/2]

- Narrow dependency**
  - Pipelined execution on one cluster node
  - e.g. a map followed by a filter
  - Failure recovery is more straightforward
- Wide dependency**
  - Requires data from all parent partitions to be available and to be shuffled across the nodes
  - Failure recovery could involve a large number of RDDs
    - Complete re-execution may be required

9/8/2016 CS535 Big Data - Fall 2016 W3.B.33

## Interface used to represent RDDs in Spark

- partitions()**
  - Returns a list of partition objects
- preferredLocations(p)**
  - List nodes where partition p can be accessed faster due to data locality
- dependencies()**
  - Return a list of dependencies
- iterator(p, parentIters)**
  - Compute the elements of partition p given iterators for its parent partitions
- partitioner()**
  - Return metadata specifying whether the RDD is hash/range partitioned

9/8/2016 CS535 Big Data - Fall 2016 W3.B.34

## In-Memory Cluster Computing: Apache Spark Spark Cluster

9/8/2016 CS535 Big Data - Fall 2016 W3.B.35

## Spark cluster and resources

The diagram illustrates the Spark cluster architecture:
 

- Driver program:** Contains the **SparkContext**.
- Cluster Manager:** Manages the cluster and receives requests from the Driver program.
- Executors:** Run on the cluster nodes. Each executor contains **Tasks** and a **Cache**.
- Backend:** The cluster can run on **Hadoop YARN**, **Mesos**, or **Standalone**.

9/8/2016 CS535 Big Data - Fall 2016 W3.B.36

## Spark cluster [1/3]

- Each application gets its own executor processes
  - Must be up and running for the duration of the entire application
- Run tasks in multiple threads
- Isolate applications from each other
  - Scheduling side (each driver schedules its own tasks)
  - Executor side (tasks from different applications run in different JVMs)
- Data cannot be shared across different Spark applications (instances of SparkContext) without writing it to an external storage system

9/8/2016 CS535 Big Data - Fall 2016 W3.B.37

## Spark cluster [2/3]

- Spark is agnostic to the underlying cluster manager
  - As long as it can acquire executor processes, and these communicate with each other, it is relatively easy to run it even on a cluster manager that also supports other applications (e.g. Mesos/YARN)

9/8/2016 CS535 Big Data - Fall 2016 W3.B.38

## Spark cluster [3/3]

- Driver program must listen for and accept incoming connections from its executors throughout its lifetime
  - Driver program must be network addressable from the worker nodes
- Driver program should run close to the worker nodes
  - On the same local area network

9/8/2016 CS535 Big Data - Fall 2016 W3.B.39

## Cluster Manager Types

- Standalone
  - Simple cluster manager included with Spark
- Mesos
  - Fine-grained sharing option
    - Frequently shared objects for Interactive applications
  - Mesos master determines the machines that handle the tasks
- Hadoop YARN
  - Resource manager in Hadoop 2

9/8/2016 CS535 Big Data - Fall 2016 W3.B.40

## Dynamic Resource Allocation

- Dynamically adjust the resources that the applications occupy
  - Based on the workload
  - Your application may give resources back to the cluster if they are no longer used
- Only available on coarse-grained cluster managers
  - Standalone mode, YARN mode, Mesos coarse grained mode

9/8/2016 CS535 Big Data - Fall 2016 W3.B.41

## In-Memory Cluster Computing: Apache Spark Scheduling

9/8/2016 CS535 Big Data - Fall 2016 W3.B.42

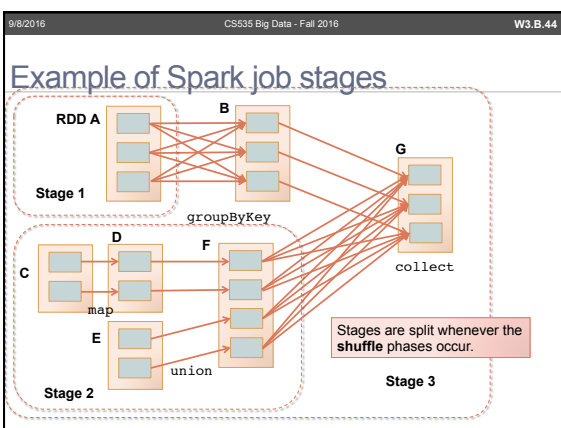
## Jobs in Spark application

- “Job”
  - A Spark action (e.g. save, collect) and any tasks that need to run to evaluate that action
- Within a given Spark application, multiple parallel tasks can run simultaneously
  - If they were submitted from separate threads

9/8/2016 CS535 Big Data - Fall 2016 W3.B.43

## Job scheduling

- User runs an action (e.g. count or save) on an RDD
- Scheduler examines that RDD's lineage graph to build a DAG of stages to execute
- Each stage contains as many pipelined transformations as possible
  - With narrow dependencies
- The boundaries of the stages are the shuffle operations
  - For wide dependencies
  - For any already computed partitions that can short circuit the computation of a parent RDD



9/8/2016 CS535 Big Data - Fall 2016 W3.B.45

## Default FIFO scheduler

- By default, Spark's scheduler runs jobs in **FIFO** fashion
- First job gets the first priority on all available resources
  - Then the second job gets the priority, etc.
  - As long as the resource is available, jobs in the queue will start right away

9/8/2016 CS535 Big Data - Fall 2016 W3.B.46

## Fair Scheduler

- Assigns tasks between jobs in a “**round robin**” fashion
  - All jobs get a roughly equal share of cluster resources
- Short jobs that were submitted when a long job is running can start receiving resources right away
  - Good response times, without waiting for the long job to finish
- Best for multi-user settings

9/8/2016 CS535 Big Data - Fall 2016 W3.B.47

## Fair Scheduler Pools

- Supports grouping jobs into pools
  - With different options (e.g. weights)
  - “high-priority” pool for more important jobs
- This approach is modeled after the **Hadoop Fair Scheduler**
- **Default behavior of pools**
  - Each pool gets **an equal share of the cluster**
  - Inside each pool, jobs run in **FIFO** order
  - If the Spark cluster creates one pool per user
    - Each user will get an equal share of the cluster
    - Each user's queries will run in order



