

CS535 BIG DATA

PART 1. BATCH COMPUTING MODELS FOR BIG DATA ANALYTICS
2. LARGE SCALE DATA ANALYSIS USING SPARK WITH CASE STUDY

Sangmi Lee Pallickara
Computer Science, Colorado State University
<http://www.cs.colostate.edu/~cs535>

9/29/16 CS535 Big Data - Fall 2016 W6.B.1

FAQs

- Sign up sheet for PA1 is available
- Term Project
 - Google computing cluster credit is available
 - Optional

9/29/16 CS535 Big Data - Fall 2016 W6.B.2

Objectives

- Large scale data analysis using Spark with case study
 - Decision tree/Random Forest
- Recommendation systems
 - Collaborative filtering
 - Latent factor approach

9/29/16 CS535 Big Data - Fall 2016 W6.B.3

Large scale data analysis using Spark with case study
Predicting Forest Cover with Decision Trees

9/29/16 CS535 Big Data - Fall 2016 W6.B.4

Tuning Decision Trees

- Spark tries a number of combinations of impurity measure, maximum depth or number of bins and reports the results

```
val evaluations =
  for (impurity <- Array("gini", "entropy");
    depth <- Array(1, 20);
    bins <- Array(10, 300))
  yield {
    val model = DecisionTree.trainClassifier(trainData, 7,
      Map[Int, Int](), impurity, depth, bins)
    val predictionsAndLabels = cvData.map(example =>
      (model.predict( example.features), example.label) )
    val accuracy =
      new MulticlassMetrics(predictionsAndLabels).
        precision ((impurity, depth, bins), accuracy)
  }
evaluations.sortBy(_._2). reverse.foreach( println) ...
```

9/29/16 CS535 Big Data - Fall 2016 W6.B.5

Tuning Decision Trees

- continued

```
(( entropy, 20,300), 0.9125545571245186)
(( gini, 20,300), 0.9042533162173727)
(( gini, 20,10), 0.8854428754813863)
(( entropy, 20,10), 0.8848951647411211)
(( gini, 1,300), 0.6358065896448438)
(( gini, 1,10), 0.6355669661959777)
(( entropy, 1,300), 0.4861446298673513)
(( entropy, 1,10), 0.4861446298673513)
```

9/29/16 CS535 Big Data - Fall 2016 W6.B.6

Categorical Features Revisited

- `Map[Int, Int]()`
 - Keys
 - **Indices of features** in the input Vector
 - Values
 - Distinct **value counts**
- Empty `Map()`
 - No features should be treated as categorical
 - All are numeric
- Numeric representation of categorical features
 - It can cause errors
 - The algorithm would be trying to learn from an ordering that has no meaning

9/29/16 CS535 Big Data - Fall 2016 W6.B.7

Treating the categorical features with **one-hot** encoding

- Encodes the categorical features as several binary 0/1 values
- Any decision rule on the "numeric" features will choose thresholds between 0 and 1
 - All are equivalent since all values are 0 or 1
- Considers the values of the underlying categorical feature **individually**
 - Increases memory usage

9/29/16 CS535 Big Data - Fall 2016 W6.B.8

Converting one-hot encoding to 1-n encoding [1/3]

```
val data = rawData.map { line =>
  val values = line.split(',').map(_.toDouble)
  val wilderness = values.slice(10, 14).indexOf(1.0).toDouble
  val soil = values.slice(14, 54).indexOf(1.0).toDouble
  val featureVector = Vectors.dense(values.slice(0, 10) ++ wilderness ++ soil)
  val label = values.last - 1
  LabeledPoint(label, featureVector)
}
```

- 4 "wilderness" features
- 40 "soil" features
- Add derived features back to first 10

9/29/16 CS535 Big Data - Fall 2016 W6.B.9

Converting one-hot encoding to 1-n encoding [2/3]

```
val evaluations =
  for (impurity <- Array("gini", "entropy"); depth <-
    Array(10, 20, 30); bins <- Array(40, 300))
  yield {
    val model =
      DecisionTree.trainClassifier(trainData, 7, Map(10->4, 11->40),
        impurity, depth, bins)
    val trainAccuracy = getMetrics(model, trainData).
    precision val cvAccuracy = getMetrics(model, cvData).
    precision ((impurity, depth, bins), (trainAccuracy, cvAccuracy))
  }
```

- Specify value count for categorical features 10, 11
 - Causes these features to be treated as categorical

9/29/16 CS535 Big Data - Fall 2016 W6.B.10

Converting one-hot encoding to 1-n encoding [3/3]

```
((entropy, 30, 300), (0.9996922984231909, 0.9438383977425239))
((entropy, 30, 40), (0.9994469978654548, 0.938934581368939))
((gini, 30, 300), (0.9998622874061833, 0.937127912178671))
((gini, 30, 40), (0.9995180059216415, 0.9329467634811934))
((entropy, 20, 40), (0.9725865867933623, 0.9280773598540899))
((gini, 20, 300), (0.9702347139020864, 0.9249630062975326))
((entropy, 20, 300), (0.9643948392205467, 0.9231391307340239))
((gini, 20, 40), (0.9679344832334917, 0.9223820503114354))
((gini, 10, 300), (0.7953203539213661, 0.7946763481193434))
((gini, 10, 40), (0.7880624698753701, 0.7860215423792973))
...
```

- Tree-building process completes several times **faster**
- By treating categorical features as categorical features, it **improves accuracy by almost 3%**

9/29/16 CS535 Big Data - Fall 2016 W6.B.11

Does decision tree algorithm build the same tree every time?

- Over N values
 - There are 2^{N-1} possible decision rules
- Decision trees use several heuristics to narrow down the rules to be considered
 - The process of picking rules involves some randomness
 - Only a few features, picked at random, are looked at each time
 - Only values from a random subset of the training data are looked
 - Trades a bit of **accuracy** for a lot of **speed**
- Decision tree algorithm **won't build the same tree every time**

9/29/16 CS535 Big Data - Fall 2016 W6.B.12

RandomForest

```
val forest = RandomForest.trainClassifier(
  trainData, 7, Map( 10 -> 4, 11 -> 40), 20,
  "auto", "entropy", 30, 300)
```

- Number of trees to build
 - Here 20
- "auto"
 - The strategy for choosing which features to evaluate at each level of the tree
 - The random decision forest implementation will NOT even consider every feature as the basis of a decision rule
 - Only a subset of all features

9/29/16 CS535 Big Data - Fall 2016 W6.B.13

Making predictions

- The results of the DecisionTree and RandomForest training
 - DecisionTreeModel and RandomForestModel objects
- predict() method
 - Accepts a Vector object
- We can classify a new example by converting it to a feature vector in the same way and predicting its target class

```
val input = "2709,125,28,67,23,3224,253,207,61,6094,0,29"
val vector = Vectors.dense(input.split(',').map(_.toDouble))
forest.predict(vector)
```

9/29/16 CS535 Big Data - Fall 2016 W6.B.14

Large scale data analysis using Spark

CASE STUDY: Recommending Music and the Audioscrobbler Data Set

9/29/16 CS535 Big Data - Fall 2016 W6.B.15

This material is built based on

- Yehuda Koren, Robert Bell, and Chris Volinsky. 2009. Matrix Factorization Techniques for Recommender Systems. *Computer* 42, 8 (August 2009), 30-37. DOI=10.1109/MC.2009.263 <http://dx.doi.org/10.1109/MC.2009.263>
- Yifan Hu, Yehuda Koren, and Chris Volinsky. 2008. Collaborative Filtering for Implicit Feedback Datasets. In *Proceedings of the 2008 Eighth IEEE International Conference on Data Mining (ICDM '08)*. IEEE Computer Society, Washington, DC, USA, 263-272. DOI=http://dx.doi.org/10.1109/ICDM.2008.22
- Sandy Ryza, Uri Laserson, Sean Owen, and Josh Wills, *Advanced Analytics with Spark*, O'Reilly, 2015


9/29/16 CS535 Big Data - Fall 2016 W6.B.16

"What percentage of the top 10,000 titles in any online media store (Netflix, iTunes, Amazon, or any other) will rent or sell at least once a month?"

9/29/16 CS535 Big Data - Fall 2016 W6.B.17

The long tail phenomenon [1/2]

- Distribution of numbers with a portion that has a large number of occurrences far from the "head" or central part of the distribution
 - The vertical axis represents popularity
 - The items are ordered on the horizontal axis according to their popularity
 - The long-tail phenomenon forces online institutions to recommend items to individual users




Erik Brynjolfsson, Yu (Jeffrey) Hu, and Duncan Simester. 2011. Goodbye Pareto Principle, Hello Long Tail: The Effect of Search Costs on the Concentration of Product Sales. *Manage. Sci.* 57, 8 (August 2011), 1373-1386. DOI=http://dx.doi.org/10.1287/mnsc.1110.1371


9/29/16 CS535 Big Data - Fall 2016 W6.B.18

The long tail phenomenon [2/2]

- "Touching the Void", Joi Simpson, 1988



- "Into Thin Air: A Personal Account of the Mt. Everest Disaster", Jon Krakauer, 1997



9/29/16 CS535 Big Data - Fall 2016 W6.B.19

Recommendation systems

- Seek to predict the "rating" or "preference" that a user would give to an item

9/29/16 CS535 Big Data - Fall 2016 W6.B.20

Applications of Recommendation Systems

- Product recommendations
 - Amazon or similar online vendors
- Movie recommendations
 - Netflix offers its customers recommendations of movies they might like
- News articles
 - News services have attempted to identify articles of interest to readers based on the articles that they have read in the past
 - Blogs, YouTube

9/29/16 CS535 Big Data - Fall 2016 W6.B.21

Netflix Prize

- The Netflix Prize challenge concerned recommender systems for movies (October, 2006)
- Netflix released a training set consisting of data from almost 500,000 customers and their ratings on 18,000 movies.
 - More than 100 million ratings
- The task was to use these data to build a model to predict ratings for a hold-out set of 3 million ratings

9/29/16 CS535 Big Data - Fall 2016 W6.B.22

Large scale data analysis using Spark

CASE STUDY: Recommending Music and the Audioscrobbler Dataset

Dataset

9/29/16 CS535 Big Data - Fall 2016 W6.B.23

Dataset

- Audioscrobbler dataset
 - 2002, Richard Jones
 - Collecting and analyzing user's songs to generate recommendation
 - Started with support for Winamp and XMMS
 - iTunes, Winamp, Windows Media Player, Foobar, iPod, AmaroK, Rhythmbox, mpd, Xbox media center, Slimserver, Jinzora, mpg321, Muine, Rhapsody, YME, Soundbridge, VLC...

last.fm

9/29/16 CS535 Big Data - Fall 2016 W6.B.24

Dataset

- Confined rating system
 - "Bob **rates** Coldplay 3.5 stars."
 - Users rate music far less frequently than they play music
- Audioscrobbler dataset
 - "Bob **played** Coldplay track"
 - Each individual data carries less information
- Implicit feedback
 - User-artist connections are implied as a side effect of other actions

9/29/16 CS535 Big Data - Fall 2016 W6.B.25

Dataset

- 141,000 unique users
- 1.6 million unique artists
- 24.2 million user's plays of artist are recorded
 - User_artist_data.txt
 - http://www-etud.iro.umontreal.ca/~bergstri/audioscrobbler_data.html
- On average, each user has played songs from about 171 artists (out of 1.6 M)
 - Extremely sparse dataset

9/29/16 CS535 Big Data - Fall 2016 W6.B.26

Large scale data analysis using Spark
CASE STUDY: Recommending Music and
the Audioscrobbler Dataset
Collaborative Filtering

9/29/16 CS535 Big Data - Fall 2016 W6.B.27

Collaborative filtering [1/2]

- Collects and analyzes a large amount of information on users' behaviors, activities or preferences and predicts what users will like based on their similarity to other users
- Explicit data collection
 - Rate an item
 - Search history
 - Favorite item
 - Wish list
- Implicit data collection
 - Viewing times
 - Tracking online purchases
 - Analyzing the user's social network

9/29/16 CS535 Big Data - Fall 2016 W6.B.28

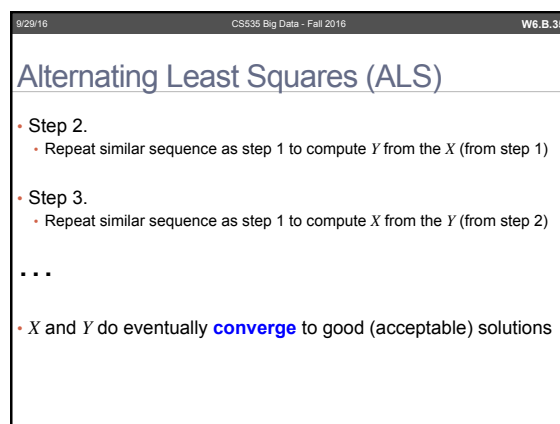
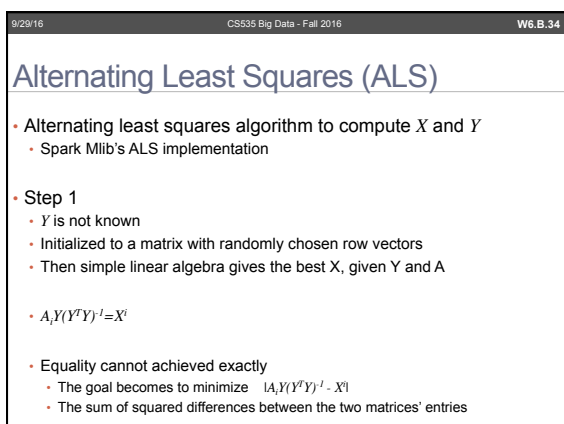
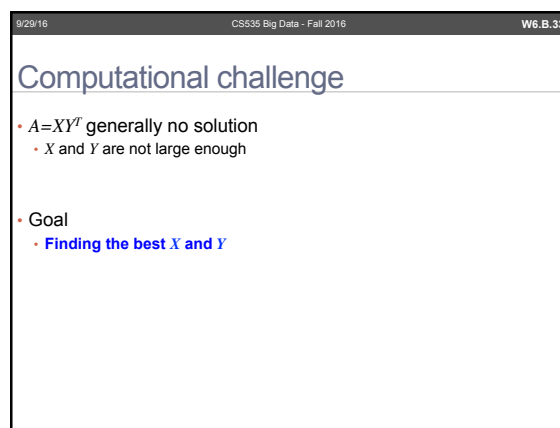
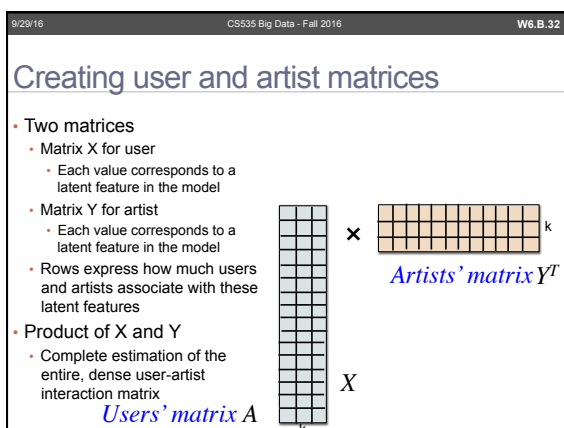
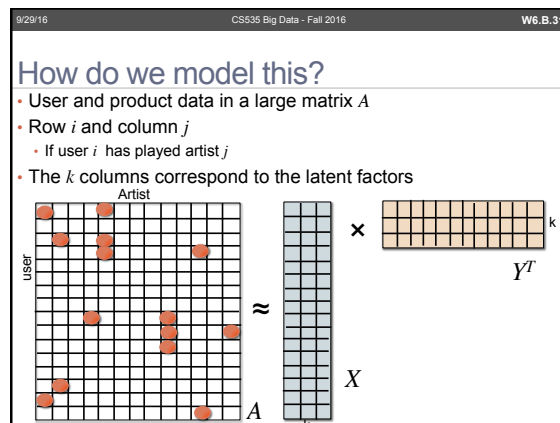
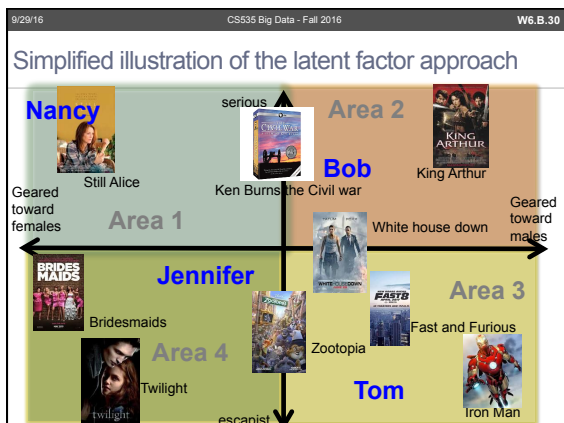
Collaborative filtering [2/2]

- Two users may share similar tastes because they are the same age
 - It is **NOT** an example of collaborative filtering
- Two users may both like the same song because they play many other same songs
 - It **IS** an example of collaborative filtering
- Algorithm that learns without access to user or artist attributes

9/29/16 CS535 Big Data - Fall 2016 W6.B.29

Latent-Factor model (1/2)

- Tries to explain **observed interactions** between large numbers of users and products through a relatively small number of **unobserved, underlying reasons**
- Within the music business context,
 - Why **millions of people** buy a particular **few of thousands of possible albums** by describing users and albums for **tens of genres and tastes that are not directly observable**



9/29/16 CS535 Big Data - Fall 2016 W6.B.36

Alternating Least Squares (ALS)

- Takes advantage of the sparsity of the input data
- Easy to apply data parallelism

9/29/16 CS535 Big Data - Fall 2016 W6.B.37

Large scale data analysis using Spark

CASE STUDY: Recommending Music and the Audioscrobbler Dataset

Building a model with Spark MLlib

9/29/16 CS535 Big Data - Fall 2016 W6.B.38

Preparing the Data

- Files are available at /user/ds/
- Spark MLlib's ALS implementation
 - Requires numeric IDs for users and items
 - Nonnegative 32-bit integers
 - An ID larger than Integer.MAX_VALUE **cannot** be used

```
val rawUserArtistData =  
  sc.textFile("hdfs:///user/ds/user_artist_data.txt")  
  
rawUserArtistData.map(_.split(' ')(0).toDouble).stats()  
rawUserArtistData.map(_.split(' ')(1).toDouble).stats()  
  
Maximum user IDs: 24443548  
Maximum artist IDs: 2147483647  
No additional transformation will be needed
```

9/29/16 CS535 Big Data - Fall 2016 W6.B.39

Extracting names

- `artist_data.txt`
- Artist ID and name separated by a tab

```
val rawArtistData =  
  sc.textFile("hdfs:///user/ds/artist_data.txt")  
  
val artistByID = rawArtistData.map { line =>  
  val (id, name) = line.split(' ')(0, 1).map(_.trim).map(_.toInt)  
}
```

- Straightforward parsing of the file into (Int, String) tuples will fail

9/29/16 CS535 Big Data - Fall 2016 W6.B.40

Extracting names

- Scala's Option class
 - Option represents a value that might only optionally exist

```
val artistByID = rawArtistData.flatMap { line =>  
  val (id, name) = line.split(' ')(0, 1).map(_.trim).map(_.toInt)  
  if (name.isEmpty) {  
    None  
  } else {  
    try {  
      Some((id, name))  
    } catch {  
      case e: NumberFormatException => None  
    }  
  }  
}
```

9/29/16 CS535 Big Data - Fall 2016 W6.B.41

Building a First Model

- Two transformations are required
 - Aliases dataset should be applied to convert all artist IDs to a canonical ID
 - The data should be converted to a Rating object
 - User-product-value data

```
import org.apache.spark.mllib.recommendation._  
  
val bArtistAlias = sc.broadcast(artistAlias)  
val trainData = rawUserArtistData.map { line =>  
  val Array(userID, artistID, count) = line.split(' ').map(_.toInt)  
  val finalArtistID = bArtistAlias.value.getOrElse(artistID, artistID)  
  Rating(userID, finalArtistID, count)  
}.cache()
```

9/29/16 CS535 Big Data - Fall 2016 W6.B.42

cache()

- RDD should be temporarily stored after being computed
- ALS is iterative
 - It will typically need to access this RDD ≥ 10 times
 - Otherwise, this RDD could be repeatedly recomputed from the original data each time

Storage Level	Cached Partitions	Fraction Cached	Size in Memory
Memory Deserialized 1x Replicated	120	100%	886.8 MB

9/29/16 CS535 Big Data - Fall 2016 W6.B.43

Broadcast variables

- For the case that many tasks (from different closures) need access to the same (immutable) data structure
- Extends normal handling of task closures
 - Caching data as raw Java objects on each executor
 - Caching data across multiple jobs and stages
- Spark will send, and hold in memory, just one copy for each executor in the cluster
 - Saves network traffic and memory

9/29/16 CS535 Big Data - Fall 2016 W6.B.44

Building the ALS model

- Constructs model as a `MatrixFactorizationModel`

```
val model = ALS.trainImplicit(trainData, 10, 5, 0.01, 1.0)
```

9/29/16 CS535 Big Data - Fall 2016 W6.B.45

Retrieving some feature vectors

- Array of 10 numbers

```
val model = ALS.trainImplicit(trainData, 10, 5, 0.01, 1.0)
model.userFeatures.mapValues(_.mkString(",")).first()

...
(4293,-0.3233030601963864, 0.31964527593541325,
0.49025505511361034, 0.09000932568001832, 0.4429537767744912,
0.4186675713407441, 0.8026858843673894, -0.4841300444834003,
-0.12485901532338621, 0.19795451025931002)
```

9/29/16 CS535 Big Data - Fall 2016 W6.B.46

Spot Checking Recommendations

- To see if the artist recommendations for user(2093760) makes any intuitive sense

```
val rawArtistsForUser =
  rawUserArtistData.map(_.split(' ')).
  filter { case Array(user,_,_) => user.toInt == 2093760 }

val existingProducts = rawArtistsForUser.map { case
  Array(_, artist,_) => artist.toInt }.collect().toSet

artistByID.filter { case (id, name) =>
  existingProducts.contains(id)
}.values.collect().foreach(println)
...
David Gray
Blackalicious
Jurassic
The Saw Doctors
Xzibit
```

9/29/16 CS535 Big Data - Fall 2016 W6.B.47

Spot Checking Recommendations

- To see five recommendations for this user (ID: 2093760)

```
val recommendations =
  model.recommendProducts(2093760, 5)
recommendations.foreach(println)

...
Rating( 2093760,1300642,0.02833118412903932)
Rating( 2093760,2814,0.027832682960168387)
Rating( 2093760,1037970,0.02726611004625264)
Rating( 2093760,1001819,0.02716011293509426)
Rating( 2093760,4605,0.027118271894797333)
```


9/29/16 CS535 Big Data - Fall 2016 W6.B.48

Large scale data analysis using Spark
CASE STUDY: Recommending Music and
the Audioscrobbler Dataset
Evaluating the Recommendation Model

9/29/16 CS535 Big Data - Fall 2016 W6.B.49

What is a “good” recommendation?

- “a popular artist”?
- “artists the user has listened to”?
- “artists the user will listen to”?

9/29/16 CS535 Big Data - Fall 2016 W6.B.50

Preparing data for evaluation

- To perform a meaningful evaluation, some of the artist play data can be set aside
 - Hidden from the ALS model building process
- The held-out data can be used as a collection of good recommendations for each user
- Compute the recommender's score

For building model	For testing model
--------------------	-------------------