

CS535 BIG DATA

PART 2. SCALABLE FRAMEWORKS FOR REAL-TIME BIG DATA ANALYTICS
1. SPEED LAYER: APACHE STORM

Sangmi Lee Pallickara
 Computer Science, Colorado State University
<http://www.cs.colostate.edu/~cs535>

10/25/2016 CS535 Big Data - Fall 2016 W10.A.1

FAQs

- Google credit available
- Assignment 2 has been posted
- URL for zookeeper has been updated

10/25/2016 CS535 Big Data - Fall 2016 W10.A.2

Today's topics

- Storm model
 - Cluster architecture
 - Trident

10/25/2016 CS535 Big Data - Fall 2016 W10.A.3

Match-making topologies and nodes

- Nimbus match-makes between the pending topologies and the Supervisor
 - Supervisor contacts Nimbus
 - Heartbeat protocol
 - Advertising the current topologies
 - Any vacancies for future topologies

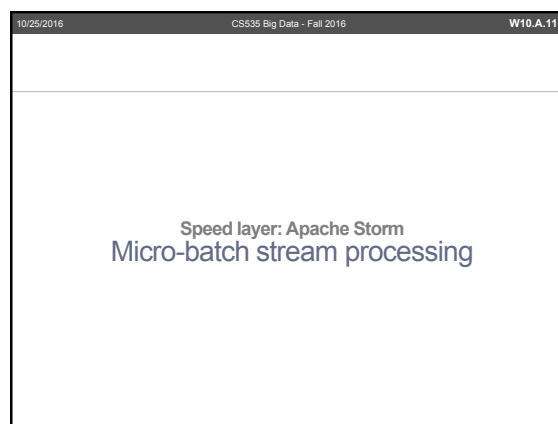
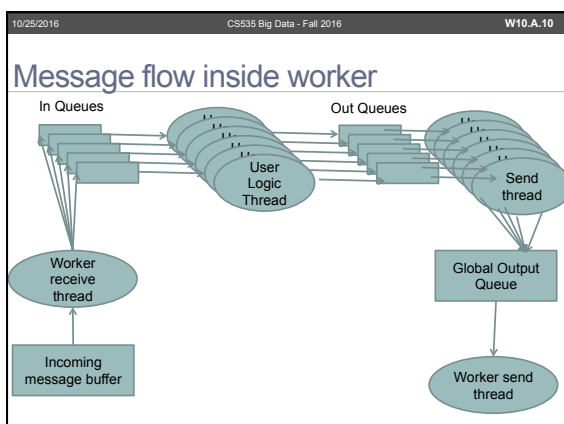
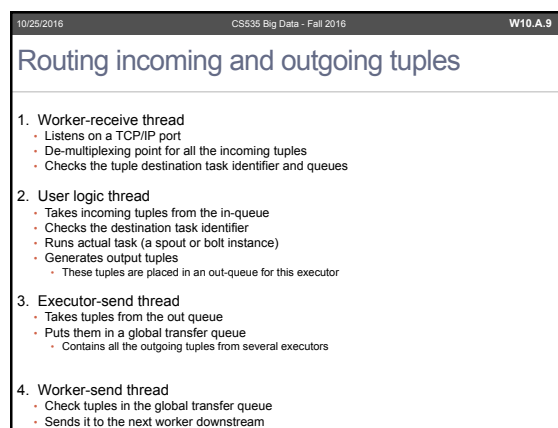
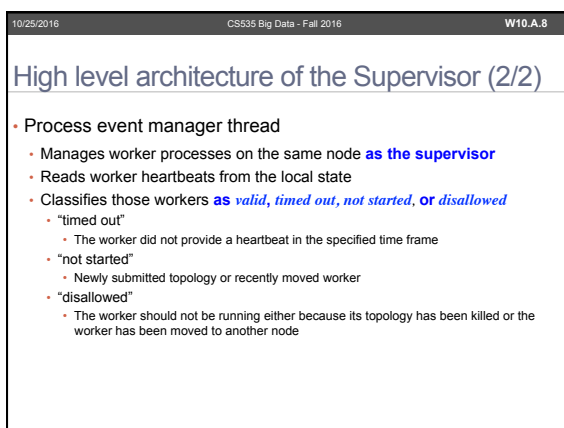
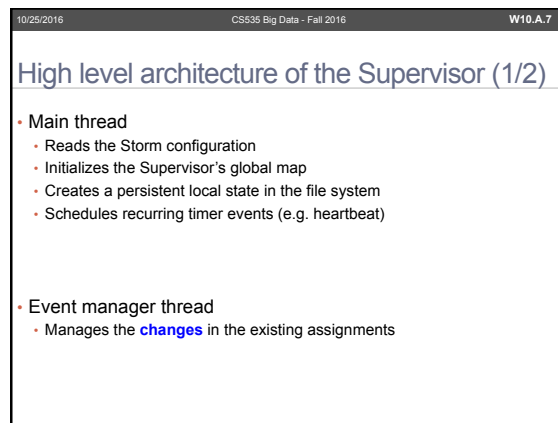
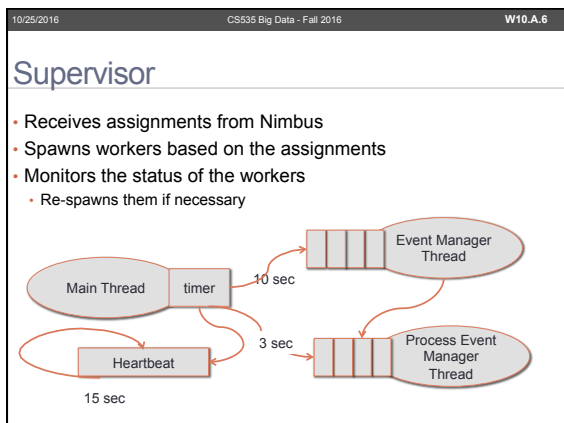
10/25/2016 CS535 Big Data - Fall 2016 W10.A.4

Coordination between Nimbus and Supervisors

- Using Zookeeper
- Nimbus and Supervisor daemons are stateless
- Their states are stored in Zookeeper or in the local disk
- If Nimbus fails,
 - Workers still continue to make forward progress
 - Users cannot submit new topologies
 - Reassigning of failed workers is not available

10/25/2016 CS535 Big Data - Fall 2016 W10.A.5

Revisit Workers/Executors/Tasks



10/25/2016 CS535 Big Data - Fall 2016 W10.A.12

Achieving exactly-once semantics

- With one-at-a-time stream processing
 - Tuples are processed independently of each other
- Micro-batch stream processing
 - Small batches of tuples are processed at one time
 - If anything in a batch fails, the entire batch is replayed
 - Batches are processed in a strict order
 - Exactly-once semantics

10/25/2016 CS535 Big Data - Fall 2016 W10.A.13

Strongly ordered processing

- If you want accuracy in your stream computing, regardless of how many failures there are:
 - Exactly once processing

```
Process(tuple){
    counter.increment()
}
```

- What if there is a failure?
 - Tuples will be replayed
 - For `counter.increment()`, you have no idea if that was processed or not

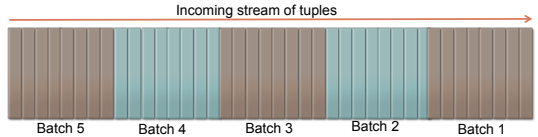
10/25/2016 CS535 Big Data - Fall 2016 W10.A.14

Exactly-once semantics

- Track ID
 - Store the ID of the latest tuple that was processed along with the count
- If the stored ID is the same as that of the current tuple ID?
 - Do nothing
- If the stored ID is different from the current tuple ID?
 - Increment the counter and update the stored ID
- You can use Ack/Nack to track tuples and maintain a queue for the tuples
 - What is the problem of this approach?

10/25/2016 CS535 Big Data - Fall 2016 W10.A.15

Micro-batch stream processing



- Batches are processed in order
 - Each batch has a unique ID
 - Always the same on every replay
- Batches must be processed to completion before moving on to the next batch

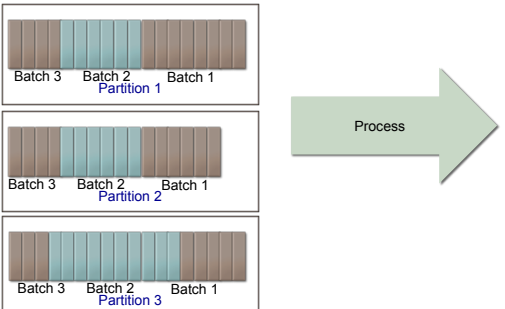
10/25/2016 CS535 Big Data - Fall 2016 W10.A.16

Micro-batch processing topologies

- Suppose that you are building a streaming application that computes the top-3 most frequently occurring words
 - Micro-batch can accomplish this task while being fully parallelized and being fault tolerant and accurate
- Task 1
 - Keeps state on the frequency of each word
 - This can be done using key/value storage
- Task 2
 - If any of the words has higher frequency than one of the current top-3 most frequent words, then the top-3 list must be updated

10/25/2016 CS535 Big Data - Fall 2016 W10.A.17

Each batch includes tuples from all partitions



10/25/2016 CS535 Big Data - Fall 2016 W10.A.18

Parallelizing the global count example

Part 1: Counting and storing the state

- The words should be re-partitioned
 - Same word is always processed by the same task** (bolt)
 - Database update is done by only one thread per-word
 - No race condition
- Stores **count and batch ID**
- For failures
 - When a failed batch is replayed:
 - If the state has current batch ID?
 - No update
 - If the state has a non-current batch ID?
 - Update

10/25/2016 CS535 Big Data - Fall 2016 W10.A.19

Parallelizing the global count example

Part 2: Computing the top-3 most frequent words

- What if we direct any new counts for every word to a single task?
 - Not scalable!
 - The single task will be a bottleneck
- What if each word counting task computes the local top-3 words and sends them to the global top-3 task?
 - Better solution

10/25/2016 CS535 Big Data - Fall 2016 W10.A.20

Failure scenario

- If a node failed and one of the top-3 lists was not sent to the global top-3 task?
 - When the batch is replayed it will be updated
- If a node failed after it updated the top-3 list
 - Update won't change the value
 - Idempotent operation

10/25/2016 CS535 Big Data - Fall 2016 W10.A.21

Speed layer: Apache Storm
 Trident Topology

10/25/2016 CS535 Big Data - Fall 2016 W10.A.22

Trident Topologies

- Trident is a Java API that translates micro-batch processing topologies into the spouts and bolts of Storm
 - Eliminates the details of transactional processing and state management
- Batching of tuples into a discrete set of transactions
- Abstracting operations on the data such as functions, filters and aggregations

10/25/2016 CS535 Big Data - Fall 2016 W10.A.23

The diagram illustrates a Trident topology. On the left, a stream of tuples (sentences) is shown. An arrow points to the right, where the tuples are grouped into three batches: Batch 1, Batch 2, and Batch 3. Each batch contains a subset of the original tuples.

10/25/2016 CS535 Big Data - Fall 2016 W10.A.24

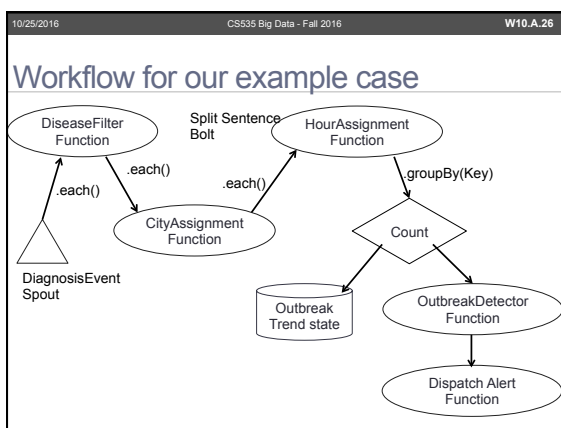
Example case (1/2)

- Collecting medical reports to identify the outbreak of a disease
- The topology will process diagnosis events that contain:
 - Latitude
 - Longitude
 - Timestamp
 - Diagnosis Code(ICD9-CM)
- E.g.
 - {39.9522, -75.1642, "03/13/2013 at 3:30 PM", "320.0 (Hemophilus meningitis)"}
- Each event includes the Global Positioning System (GPS) coordinates of the occurrence

10/25/2016 CS535 Big Data - Fall 2016 W10.A.25

Example case (2/2)

- To detect an outbreak,
 - The system will **count** the occurrence of specific disease codes within geographic location over a specified period of time
 - The system will **group** the occurrences by hour and calculate a trend against the moving average
 - The system will **use a simple threshold to determine** if there is an outbreak
 - If the count of occurrences for the hour is greater than some threshold, the system will **send an alert**



10/25/2016 CS535 Big Data - Fall 2016 W10.A.27

```

public class OutbreakDetectionTopology {
    public static StormTopology buildTopology() {
        TridentTopology topology = new TridentTopology();
        DiagnosisEventSpout spout = new DiagnosisEventSpout();
        Stream inputStream = topology.newStream("event", spout);
        inputStream
            // Filter for critical events.
            .each(new Fields("event"), new DiseaseFilter())
            // Locate the closest city
            .each(new Fields("event"),
                new CityAssignment(),
                new Fields("city"))
            // Derive the hour segment
            .each(new Fields("event", "city"),
                new HourAssignment(),
                new Fields("hour", "cityDiseaseHour"))
            // Group occurrences in same city and hour
            .groupBy(new Fields("cityDiseaseHour"))
            // Count occurrences and persist the results.
            .persistentAggregate(new OutbreakTrendFactory(),
                new Count(),
                new Fields("count"))
            .newValuesStream()
    }
}
    
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.28

continued

```

// Detect an outbreak
.each(new Fields("cityDiseaseHour", "count"),
    new OutbreakDetector(),
    new Fields("alert"))
// Dispatch the alert
.each(new Fields("alert"),
    new DispatchAlert(),
    new Fields());
}
    
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.29

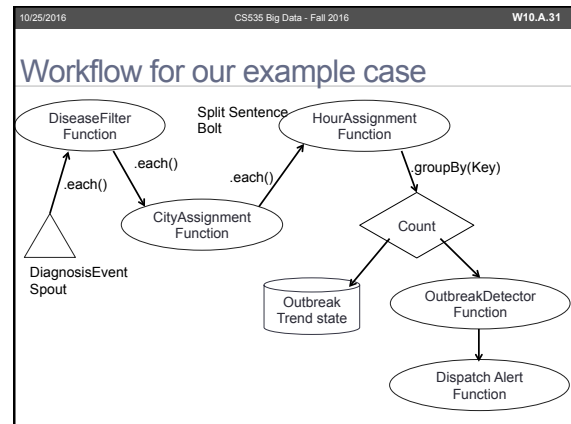
Introducing Trident Spout

- Batch**
 - Trident spouts must emit tuples in batches
- Composition of a batch**
 - Non-transactional**
 - No guarantee on the composition of the batches and might overlap
 - Two different batches might contain the same tuples
 - Transactional**
 - Guaranteed and non-overlapping
 - Same batch contains the same tuples
 - Opaque**
 - Guaranteed and non-overlapping
 - Contents of a batch may change

10/25/2016 CS535 Big Data - Fall 2016 W10.A.30

Trident Spout interface

```
public interface ITridentSpout < T > extends Serializable
{
    BatchCoordinator < T > getCoordinator(String txStateId, Map conf, TopologyContext context);
    Emitter < T > getEmitter(String txStateId, Map conf, TopologyContext context);
    Map getComponentConfiguration();
    Fields getOutputFields();
}
```



10/25/2016 CS535 Big Data - Fall 2016 W10.A.32

DiagnosisEventSpout

```
public class DiagnosisEventSpout implements ITridentSpout
< Long > {
    private static final long serialVersionUID = 1L;
    SpoutOutputCollector collector;
    BatchCoordinator < Long > coordinator = new
        DefaultCoordinator();
    Emitter < Long > emitter = new DiagnosisEventEmitter();
    @Override
    public BatchCoordinator < Long > getCoordinator(
        String txStateId, Map conf, TopologyContext
        context)
    {
        return coordinator;
    }
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.33

Continued

```

    @Override
    public Emitter < Long > getEmitter( String
    txStateId, Map conf, TopologyContext context) {
        return emitter;
    }
    @Override
    public Map getComponentConfiguration() {
        return null;
    }
    @Override
    public Fields getOutputFields() {
        return new Fields("event");
    }
}

```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.34

BatchCoordinator

```
public class DefaultCoordinator implements BatchCoordinator < Long >,
Serializable {
    private static final long serialVersionUID = 1L;
    private static final Logger LOG =
    LoggerFactory.getLogger( DefaultCoordinator.class);
    @Override
    public boolean isReady( long txid) {
        return true;
    }
    @Override
    public void close() {
    }
    @Override
    public Long initializeTransaction(long txid, Long prevMetadata) {
        LOG.info(" Initializing Transaction [" + txid + "]");
        return null;
    }
    @Override
    public void success(long txid) {
        LOG.info(" Successful Transaction [" + txid + "]");
    }
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.35

Emitter

```
public class DiagnosisEventEmitter implements Emitter < Long >, Serializable {
    private static final long serialVersionUID = 1L;
    AtomicInteger successfulTransactions = new AtomicInteger( 0);
    @Override
    public void emitBatch( TransactionAttempt tx,
        Long coordinatorMeta,
        TridentCollector collector) {
        for (int i = 0; i < 100000; i++) {
            List < Object > events =
                new ArrayList < Object > ();
            double lat = new Double(-30
                + (int) (Math.random() * 75));
            double lng = new Double(-120
                + (int) (Math.random() * 70));
            long time = System.currentTimeMillis();
            String diag = new Integer( 320
                + (int) (Math.random() * 7)).toString();
            DiagnosisEvent event =
                new DiagnosisEvent( lat, lng, time, diag);
            events.add(event);
            collector.emit(events);
        }
    }
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.36

continued

```
@Override
public void success( TransactionAttempt tx) {
    successfulTransactions.incrementAndGet();
}
@Override
public void close() {
}
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.37

DiagnosisEvent

- ICD-9-CM codes
- 320 Bacterial meningitis
- 321 Meningitis due to other organisms
- 322 Meningitis of unspecified cause
- 323 Encephalitis myelitis and encephalomyelitis
- ...

10/25/2016 CS535 Big Data - Fall 2016 W10.A.38

DiagnosisEvent class

```
public class DiagnosisEvent implements Serializable {
    private static final long serialVersionUID = 1L;
    public double lat;
    public double lng;
    public long time;
    public String diagnosisCode;
    public DiagnosisEvent(double lat, double lng, long time, String
        diagnosisCode) {
        super();
        this.time = time;
        this.lat = lat;
        this.lng = lng;
        this.diagnosisCode = diagnosisCode;
    }
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.39

Trident operations - filters and functions

- Operations
 - Adding the logic components that implement the business process
- Filters
- Functions
- Join
- Aggregation
- Group

- Implementing methods on the Stream object

10/25/2016 CS535 Big Data - Fall 2016 W10.A.40

Methods on the Stream object

```
public class Stream implements IAggregatableStream {
    public Stream each(Fields inputFields, Filter filter) {
        ...
    }
    public IAggregatableStream each(Fields inputFields,
        Function function, Fields functionFields){
        ...
    }
    public GroupedStream groupBy(Fields fields) {
        ...
    }
    public TridentState persistentAggregate(StateFactory
        stateFactory, CombinerAggregator agg, Fields
        functionFields) {
        ...
    }
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.41

Example code with the disease detecting example

```
inputStream.each(new Fields(" event"), new
    DiseaseFilter())
    .each(new Fields("event"),
        new CityAssignment(),
        new Fields("city"))
    .each(new Fields("event", "city"),
        new HourAssignment(),
        new Fields("hour", "cityDiseaseHour"))
    .groupBy(new Fields("cityDiseaseHour"))
    .persistentAggregate(new OutbreakTrendFactory(),
        new Count(), new Fields("count"))
    .newValuesStream()
    .each(new Fields("cityDiseaseHour", "count"),
        new OutbreakDetector(), new Fields("alert"))
    .each(new Fields("alert"),
        new DispatchAlert(), new Fields());
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.42

Trident filters

- For example, the system wants to ignore disease events that are not of concern
 - Focus on meningitis (code 320,321,and 322)
- Providing a `BaseFilter` class

```
public interface Filter extends EachOperation {
    boolean isKeep(TridentTuple tuple);
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.43

Trident Filter interface

```
public interface Filter extends EachOperation {
    boolean isKeep(TridentTuple tuple);
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.44

DiseaseFilter

```
public class DiseaseFilter extends BaseFilter {
    private static final long serialVersionUID = 1L;
    private static final Logger LOG = LoggerFactory.getLogger(
        DiseaseFilter.class);

    @Override
    public boolean isKeep( TridentTuple tuple) {
        DiagnosisEvent diagnosis=(DiagnosisEvent)
        tuple.getValue(0);
        Integer code=Integer.parseInt( diagnosis.diagnosisCode);
        if (code.intValue() <= 322) {
            LOG.debug(" Emitting disease [" +
                diagnosis.diagnosisCode + "]");
            return true;
        } else {
            LOG.debug(" Filtering disease [" +
                diagnosis.diagnosisCode + "]");
            return false;
        }
    }
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.45

Applying filter to each tuple

```
inputStream.each(new Fields("event"), DiseaseFilter())
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.46

Trident functions

- Consume **tuples** and optionally emit new **tuples**
- Trident functions are additive
 - The values emitted by functions are fields that are added to the tuple
 - They do not remove or mutate existing fields

```
public interface Function extends EachOperation {
    void execute(TridentTuple tuple,
        TridentCollector collector);
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.47

Writing your BaseFunction

```
public class CityAssignment extends BaseFunction {
    private static final long serialVersionUID = 1L;
    private static final Logger LOG=
        LoggerFactory.getLogger( CityAssignment.class);
    private static Map < String, double[] > CITIES = new HashMap <
        String, double[] >();

    {
        // Initialize the cities we care about.
        double[] phl = {39.875365, -75.249524 };
        CITIES.put("PHL", phl);
        double[] nyc = {40.71448, -74.00598 };
        CITIES.put("NYC", nyc);
        double[] sf = {-31.4250142, -62.0841809 };
        CITIES.put("SF", sf);
        double[] la = {-34.05374, -118.24307 };
        CITIES.put("LA", la);
    }
}
```


10/25/2016 CS535 Big Data - Fall 2016 W10.A.48

Writing your BaseFunction

```
#Override
public void execute( TridentTuple tuple, TridentCollector
collector) {
    DiagnosisEvent diagnosis=
        (DiagnosisEvent)tuple.getValue(0);
    double leastDistance = Double.MAX_VALUE;
    String closestCity = "NONE";
    // Find the closest city.
    for (Entry<String, double[]> city : CITIES.entrySet()) {
        double R = 6371;
        // km
        double x = (city.getValue()[0] - diagnosis.lng) *
            Math.cos((city.getValue()[0] + diagnosis.lng) / 2);
        double y = (city.getValue()[1] - diagnosis.lat);
        double d = Math.sqrt( x * x + y * y) * R;
        if (d < leastDistance) {
            leastDistance = d;
            closestCity = city.getKey();
        }
    }
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.49

Writing your BaseFunction

```
// Emit the value.
List<Object> values = new ArrayList<Object>();
Values.add(closestCity);
LOG.debug("Closest city to lat =[" + diagnosis.lat + "],
    lng =[" + diagnosis.lng + "] = [" + closestCity + "],
    d =[" + leastDistance + "]);
collector.emit(values);
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.50

Trident aggregator

- Allows topologies to combine tuples
 - They replace tuple fields and values
 - Function does not change
- CombinerAggregator
- ReducerAggregator
- Aggregator

10/25/2016 CS535 Big Data - Fall 2016 W10.A.51

CombinerAggregator

- Combines a set of tuples into a single field
- Storm calls the `init()` method with each tuple then repeatedly calls `combine()` method until the partition is processed

```
public interface CombinerAggregator {
    T init(TridentTuple tuple);
    T combine(T val1, T val2);
    T zero(); //emits and returns value
}
```

10/25/2016 CS535 Big Data - Fall 2016 W10.A.52

ReducerAggregator

```
public interface ReducerAggregator<T> extends
Serializable {
    T init();
    T reduce(T curr, TridentTuple tuple);
}
```

- Storm calls the `init()` method to retrieve the initial value
- Then `reduce()` is called with each tuple until the partition is fully processed
- The first parameter into the `reduce()` method is **the cumulative partial aggregation**
- The implementation should return the result of incorporating the tuple into that partial aggregation

10/25/2016 CS535 Big Data - Fall 2016 W10.A.53

Aggregator

- The most general aggregation operation

```
public interface Aggregator<T> extends Operation {
    T init(Object batchId, TridentCollector collector);
    void aggregate(T val, TridentTuple tuple,
        TridentCollector collector);
    void complete(T val, TridentCollector collector);
}
```

- The `aggregate()` method is similar to the `execute()` method of a Function interface
 - It also includes a parameter for the value
 - This allows the Aggregator to accumulate a value as it processes the tuples. Notice that with an Aggregator, the collector is passed into both the `aggregate()` method as well as the `complete()` method
 - You can emit any arbitrary number of tuples

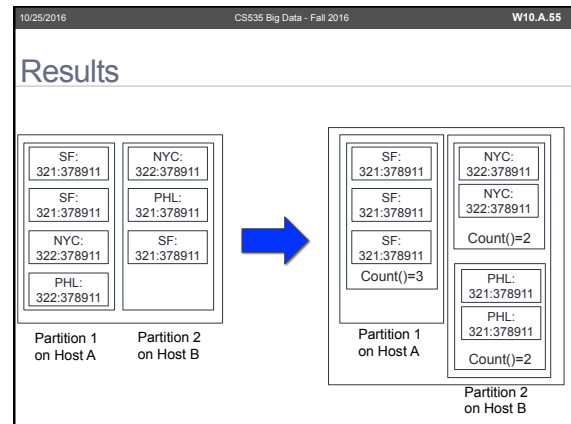
10/25/2016 CS535 Big Data - Fall 2016 W10.A.54

Writing and applying Count

```
public class Count implements CombinerAggregator < Long > {
    @Override
    public Long init( TridentTuple tuple){
        return 1L;
    }
    @Override
    public Long combine( Long val1, Long val2) {
        return val1 + val2;
    }
    @Override
    public Long zero() {
        return 0L;
    }
}
```

*Applying grouping and counting

```
.groupBy(new Fields("cityDiseaseHour"))
.persistentAggregate(new OutbreakTrendFactory(), new Count(), new
Fields(" count")). newValuesStream()
```



10/25/2016 CS535 Big Data - Fall 2016 W10.A.56

Trident state

- Trident has a first-level primitive for state
- State interface

```
public interface State {
    void beginCommit( Long transactionId);
    void commit( Long transactionId);
}
```

- Each batch (of tuples) has its own transaction identifier
- State object specifies when the state is being committed and when the commit should complete